# CS346 Redbase Part 5 - EX Component Proposal
# Distributed RedBase

Aditya Bhandari (adityasb)

## 1 Idea

The high-level idea is to make a **Distributed RedBase** by horizontal fragmentation of data across multiple nodes. The fragmentation of data and handling of queries will be done by a centralized control or *"master"* node, whereas all the data will reside on children or *"data"* nodes. All the user commands and queries will be given to the *master* node, who will then redirect the command or query to the appropriate *data* node.

This proposal document tries to list the various goals of the project extension, followed by a rough idea of the proposed system design and implementation details, finally followed by a set of milestones in order to track the progress of the project.

### 1.1 Goals

The primary goal of the extension is to create a distributed database system, consisting of data fragmented and stored across multiple nodes. The user should be able to perform the required queries as if the data is stored in a single node database. Underneath, the query will be distributed to the appropriate nodes and the results combined at the *master* node.

The aim is to be able to support all the queries from the QL component of the project in the distributed scenario as well - `INSERT`, `DELETE`, `UPDATE` and `SELECT`, in addition to the other RQL commands from the SM component - `create/destroy table`, `create/destroy index`, `load table` and so on. The user should be able to specify the distributed nature of the database through the `dbcreate` command, and specify the partitioning while creating a relation in the database.

### 1.2 Functionality

As mentioned in the previous subsection, the main functionality will be the option for the user to create a distributed database when using the `dbcreate` command - specifying the required number of *data* nodes. In the database, the RQL command to `create table` will enable the user to specify the partitioning for that relation across the *data* nodes. Similarly, the other RQL commands on the database - `help, print` should work as before even in the distributed scenario.

The functionality goal is to support all the basic RQL queries on a distributed database. The `INSERT` command should insert the required tuple in the corresponding *data* node, based on the partitioning vector. The `DELETE` command should delete the tuple from the corresponding *data* node, whereas the `UPDATE` command should perform the update at the corresponding *data* node. The `SELECT` command is the most challenging, which will perform a distributed join of the required relations by moving data temporarily across the *data* nodes.

(The movement of the data across the nodes, the communication of the *master* node with the *data* nodes and the *data* nodes among themselves will be handled by a newly added special communication layer underneath the RedBase nodes.)

# 2 System Design

The high-level design of the Distributed RedBase system is shown in the following figure and explained in the following two subsections.
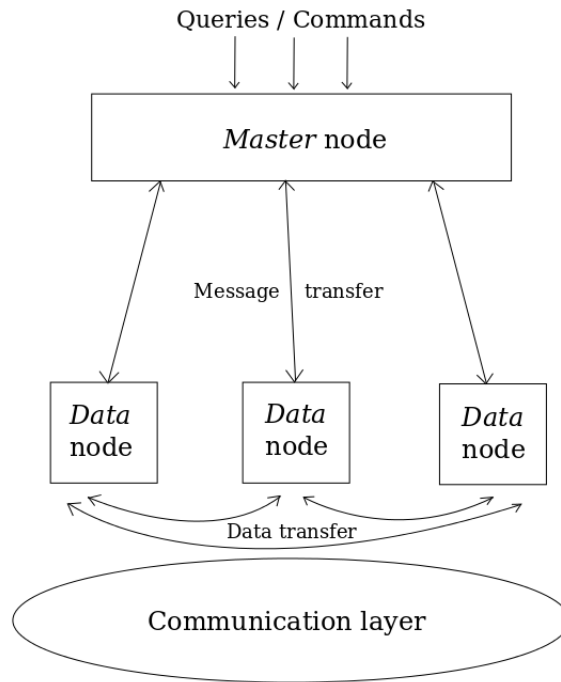


Figure 1: Distributed RedBase System

## 2.1 Distributed database - *Master* node and *Data* nodes

As explained in the previous sections, the RedBase system will consist of a *master* node to handle and control all the features and aspects of the distributed database. On the other hand, the actual data in the database will be stored in children *data* nodes, which are created by the *master* node when a distributed database is created, according to the requirements specified by the user.

When the *master* node creates the *data* nodes, it stores the partitioning information for each relation in the database in a **partitioning vector**. This vector will basically represents the ranges of the fragmentation of the relation among the different nodes. We will only consider a range partitioning scheme over any one of the attributes in the relation (for simplicity). One such vector will need to be stored in the *master* node for every relation in the database.

All the RQL queries and commands to the system are input to the *master* node. It then allocates the query to the appropriate node based on the partitioning vector for the corresponding relation. In the case of a local query - INSERT, DELETE, UPDATE, the *data* node performs the query and returns the result to the *master* node. But, in the case of a distributed query, that is, a SELECT requiring a distributed join, an entire relation will be transported to each of the *data* nodes, which will then perform a join with the local fragment of the other relation and return the result set to the *master* node. The *master* node will then evaluate the UNION of these result sets and get the final result of the query.

## 2.2 Network simulation - Communication layer

TODO

# 3 Implementation Overview

TODO

## 3.1 SM Component extension

TODO

## 3.2 QL Component extension

TODO

## 3.3 Network / Communication layer

TODO

# 4 Milestones

For the purpose of tracking progress, it makes sense to divide the EX component of the project into three milestones, each progressively more challenging than the previous one. The goals of each of these milestones are listed and explained in brief below.

## 4.1 Basic - Distributed database with multiple nodes and whole relation scans

- `dbcreate` command to create a distributed database with multiple *data* nodes

- Fragment data horizontally across the nodes (simple user-specified range partitioning)

- Communication layer for the interaction among the *data* nodes and with the *master* node

- Scan the whole relation for the `print` command (by `UNION` of the fragmented result sets)

## 4.2 Intermediate - `INSERT, UPDATE` and `DELETE`

- Issue the query from the *master* node to the appropriate *data* node (based on the partitioning vector for the corresponding relation)

- Execute the query locally at the *data* node and result returned to the *master* node

## 4.3 Ambitious - `SELECT` (Distributed join)

- Copy an entire relation temporarily at all nodes using the communication layer (Add new operators in the query plan for this step)

- Perform join locally at each *data* node and return the result set to the *master* node.

- Evaluate the final result at the *master* by evaluating the `UNION` of the result sets